

Programing Assignment 1: Resnet

Bingnan Li
2020533092

libn@shanghaitech.edu.cn

1. Resnet on CIFAR10

In this section, I will introduce the implementation details and comparison between training from scratch and training with fine-tuning.

1.1. Network Architecture

For this classification task, I use Resnet34 as the backbone and modified two layers to fit the input image size and the number of classes. The detailed modifications are as follows:

- Change the first convolution layer from 7×7 kernel with stride 2 to 3×3 kernel with stride 1.
- Change the final fully-connected layer from 1000 output to 10 output.

1.2. Global Training Settings

1.2.1 Dataset split

The original CIFAR10 dataset contains 60000 images with size 32×32 with 50000 for training and 10000 for testing.

Given that it doesn't naturally split the validation set, I randomly choose 5000 images from the training set as the validation set with a fixed seed.

1.2.2 Data Augmentation

In order to increase the robustness and generalization of the model, I use random horizontal flip and random crop with zero padding as augmentation methods.

1.3. Training from scratch

1.3.1 Training details

The hyperparameters settings follow the detail of [1], which are listed as follows:

- learning rate: 0.1
- learning rate decay: decay 0.1 after 32k and 48k iterations.




Figure 1. Training curve of Resnet34 on CIFAR10

- batch size: 128

Moreover, I used *Cross Entropy loss* as loss function and used classification accuracy as metric

1.3.2 Experimental results

After 200 epoch training, the average classification accuracy is **91.31%** under test set. The training curve are shown as Figure1

1.4. Training from Fine-tuning

Since I used *Resnet34* as backbone, I can directly load the pretrained weight of *Resnet34* trained on *ImageNet*.

1.4.1 Training Policy

Given that I only changed the first and the last of layer of *Resnet34*, I choose *semi-freezing training* policy to fine-tune.

In detail, I freeze all the middle layers except the first and the last layer, and only train those two layers with learning rate 0.01 for the first 50 epoch. Then I unfreeze all the layers and train the whole network with for the next 150 epoch. The learning rate decay changed to decay 0.1 after 10k and 35k iterations.

1.5. The ruler

The L^AT_EX style defines a printed ruler which should be present in the version submitted for review. The ruler is provided in order that reviewers may comment on particular lines in the paper without circumlocution. If you are preparing a document using a non-L^AT_EX document preparation system, please arrange for an equivalent ruler to appear on the final output pages. The presence or absence of the ruler should not change the appearance of any other content on the page. The camera-ready copy should not contain a ruler. (L^AT_EX users may use options of cvpr.sty to switch between different versions.)

Reviewers: note that the ruler measurements do not align well with lines in the paper — this turns out to be very difficult to do well when the paper contains many figures and equations, and, when done, looks ugly. Just use fractional references (e.g., this line is 087.5), although in most cases one would expect that the approximate location will be adequate.

1.6. Paper ID

Make sure that the Paper ID from the submission system is visible in the version submitted for review (replacing the “*****” you see in this document). If you are using the L^AT_EX template, make sure to update paper ID in the appropriate place in the tex file.

1.7. Mathematics

Please number all of your sections and displayed equations as in these examples:

$$E = m \cdot c^2 \tag{1}$$

and

$$v = a \cdot t. \tag{2}$$

It is important for readers to be able to refer to any particular equation. Just because you did not refer to it in the text does not mean some future reader might not need to refer to it. It is cumbersome to have to use circumlocutions like “the equation second from the top of page 3 column 1”. (Note that the ruler will not be present in the final copy, so is not an alternative to equation numbers). All authors will benefit from reading Mermin’s description of how to write mathematics: <http://www.pamitc.org/documents/mermin.pdf>.

1.8. Miscellaneous

Compare the following:

`$conf_a$` *conf_a*
`conf_a` *conf_a*

See The T_EXbook, p165.

The space after e.g., meaning “for example”, should not be a sentence-ending space. So e.g. is correct, e.g. is not. The provided `\eg` macro takes care of this.

When citing a multi-author paper, you may save space by using “et alia”, shortened to “et al.” (not “et. al.” as “et” is a complete word). If you use the `\etal` macro provided, then you need not worry about double periods when used at the end of a sentence as in Alpher et al. However, use it only when there are three or more authors. Thus, the following is correct: “Frobnication has been trendy lately. It was introduced by Alpher [?], and subsequently developed by Alpher and Fotheringham-Smythe [?], and Alpher et al. [?].”

This is incorrect: “... subsequently developed by Alpher et al. [?] ...” because reference [?] has just two authors.

2. Formatting your paper


All text must be in a two-column format. The total allowable size of the text area is $6\frac{7}{8}$ inches (17.46 cm) wide by $8\frac{7}{8}$ inches (22.54 cm) high. Columns are to be $3\frac{1}{4}$ inches (8.25 cm) wide, with a $\frac{5}{16}$ inch (0.8 cm) space between them. The main title (on the first page) should begin 1 inch (2.54 cm) from the top edge of the page. The second and following pages should begin 1 inch (2.54 cm) from the top edge. On all pages, the bottom margin should be $1\frac{1}{8}$ inches (2.86 cm) from the bottom edge of the page for 8.5×11 -inch paper; for A4 paper, approximately $1\frac{5}{8}$ inches (4.13 cm) from the bottom edge of the page.

2.1. Margins and page numbering


All printed material, including text, illustrations, and charts, must be kept within a print area $6\frac{7}{8}$ inches (17.46 cm) wide by $8\frac{7}{8}$ inches (22.54 cm) high. Page numbers should be in the footer, centered and $\frac{3}{4}$ inches from the bottom of the page. The review version should have page numbers, yet the final version submitted as camera ready should not show any page numbers. The L^AT_EX template takes care of this when used properly.

2.2. Type style and fonts

Wherever Times is specified, Times Roman may also be used. If neither is available on your word processor, please use the font closest in appearance to Times to which you have access.



(a) An example of a subfigure.



(b) Another example of a subfigure.

Figure 2. Example of a short caption, which should be centered.

MAIN TITLE. Center the title $1\frac{3}{8}$ inches (3.49 cm) from the top edge of the first page. The title should be in Times 14-point, boldface type. Capitalize the first letter of nouns, pronouns, verbs, adjectives, and adverbs; do not capitalize articles, coordinate conjunctions, or prepositions (unless the title begins with such a word). Leave two blank lines after the title.

AUTHOR NAME(s) and AFFILIATION(s) are to be centered beneath the title and printed in Times 12-point, non-boldface type. This information is to be followed by two blank lines.

The **ABSTRACT** and **MAIN TEXT** are to be in a two-column format.

MAIN TEXT. Type main text in 10-point Times, single-spaced. Do NOT use double-spacing. All paragraphs should be indented 1 pica (approx. $\frac{1}{6}$ inch or 0.422 cm). Make sure your text is fully justified—that is, flush left and flush right. Please do not place any additional blank lines between paragraphs.

Figure and table captions should be 9-point Roman type as in ?? and Fig. 2. Short captions should be centred.

Callouts should be 9-point Helvetica, non-boldface type. Initially capitalize only the first word of section titles and first-, second-, and third-order headings.

FIRST-ORDER HEADINGS. (For example, 1. Introduction) should be Times 12-point boldface, initially capitalized, flush left, with one blank line before, and one blank line after.

SECOND-ORDER HEADINGS. (For example, 1.1. Database elements) should be Times 11-point boldface, initially capitalized, flush left, with one blank line before, and one after. If you require a third-order heading (we discourage it), use 10-point Times, boldface, initially capitalized, flush left, preceded by one blank line, followed by a period and your text on the same

line.

2.3. Footnotes

Please use footnotes¹ sparingly. Indeed, try to avoid footnotes altogether and include necessary peripheral observations in the text (within parentheses, if you prefer, as in this sentence). If you wish to use a footnote, place it at the bottom of the column on the page on which it is referenced. Use Times 8-point type, single-spaced.

2.4. Cross-references

For the benefit of author(s) and readers, please use the

`\cref{...}`

command for cross-referencing to figures, tables, equations, or sections. This will automatically insert the appropriate label alongside the cross-reference as in this example:

To see how our method outperforms previous work, please see ?? and ?. It is also possible to refer to multiple targets as once, e.g. to ?? and Fig. 2a. You may also return to Sec. 2 or look at Eq. (2).

If you do not wish to abbreviate the label, for example at the beginning of the sentence, you can use the

`\Cref{...}`

command. Here is an example:

?? is also quite important.

¹This is what a footnote looks like. It often distracts the reader from the main flow of the argument.

2.5. Illustrations, graphs, and photographs

All graphics should be centered. In L^AT_EX, avoid using the center environment for this purpose, as this adds potentially unwanted whitespace. Instead use

```
\centering
```

at the beginning of your figure. Please ensure that any point you wish to make is resolvable in a printed copy of the paper. Resize fonts in figures to match the font in the body text, and choose line widths that render effectively in print. Readers (and reviewers), even of an electronic copy, may choose to print your paper in order to read it. You cannot insist that they do otherwise, and therefore must not assume that they can zoom in to see tiny details on a graphic.

When placing figures in L^AT_EX, it's almost always best to use `\includegraphics`, and to specify the figure width as a multiple of the line width as in the example below

```
\usepackage{graphicx} ...  
\includegraphics[width=0.8\linewidth]  
  {myfile.pdf}
```

2.6. Color

Please refer to the author guidelines on the CVPR 2023 web page for a discussion of the use of color in your document.

If you use color in your plots, please keep in mind that a significant subset of reviewers and readers may have a color vision deficiency; red-green blindness is the most frequent kind. Hence avoid relying only on color as the discriminative feature in plots (such as red vs. green lines), but add a second discriminative feature to ease disambiguation.

3. Final copy

You must include your signed IEEE copyright release form when you submit your finished paper. We MUST have this form before your paper can be published in the proceedings.

Please direct any questions to the production editor in charge of these proceedings at the IEEE Computer Society Press: <https://www.computer.org/about/contact>.

References

- [1] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition, pages 770–778, 2016. 1